

Instructions:

1. Run task_2.py to generate the tensor database if not already generated.
2. Press Run All (or restart kernel and run all cells).
3. You will be prompted to provide an image ID and a k value.

Task 3: Implement a program which, given an image ID and a value “k”, returns and visualizes the most similar k images based on each of the visual model –you will select the appropriate distance/similarity measure for each feature model.
For each match, also list the corresponding distance/similarity score.

```
In [ ]: IMG_ID = int(input("Provide an Image ID in the Caltech101 dataset in range [0, 8676]."))  
  
K = int(input("Enter K, the number of similar images K to find for each feature model."))
```

```
In [ ]: # Initialize dataset, downloads into <git_root>/Code/caltech101 if not present. Untracked by git.  
  
import os  
from torchvision.datasets import Caltech101  
  
dataset = Caltech101(root=os.path.abspath(""), download=True)
```

Files already downloaded and verified

```
In [ ]: # Find the k most similar images for a given Image ID for each feature model.  
  
image, category = dataset[IMG_ID][0], dataset.categories[dataset[IMG_ID][1]]  
  
if image.mode != 'RGB':  
    print("WARNING: The provided image does not have RGB channels, the image will be converted to RGB.")  
    image = image.convert('RGB')  
  
print("Input image:\nImage ID: ", IMG_ID)  
display(image)  
print("Category:", category)
```

Input image:
Image ID: 2500



Category: airplanes

```
In [ ]: # Resize to 300x100 and partition to grid.

image300x100 = image.resize((300, 100))

from utils.utilities import partition_to_grid

grid = partition_to_grid(image300x100, num_rows=10, num_cols=10, hor_pixels=30, ver_pixels=10)
```

```
In [ ]: # Color moments (CM10x10).

# Find K most similar images by color moments.

from feature_extractors.color_moments import get_color_vector

color_vector = get_color_vector(grid)

# Load color tensor database into memory.

from utils.utilities import tensor_loader

color_tensor_dict = tensor_loader('color_tensors.pt', os.path.abspath(""))

from utils.utilities import top_k_ranker

# Manhattan distance (City-block distance): This distance can be imagined as the length needed
# to move between two points in a grid where you can only move up, down, left or right. This
# measure is sensitive to differences along each dimension of the feature vectors, capturing
# the total absolute difference between corresponding moments, which makes it great for
```

```

# assessing differences in color distribution in terms of absolute deviations. Since it
# does not square the differences like Euclidean distance, it is less susceptible outliers
# when many dimensions differ.

# We use Earth Mover's Distance (Wasserstein distance) as color moments can be thought
# of as probability distributions over color space. EMD is capable of comparing distributions
# with different shapes and scales, color moments consider the spatial arrangement of colors,
# not just the frequency of colors. EMD takes this spatial information into account by
# considering how color probability mass is transported from one distribution to another.

from scipy.spatial.distance import cityblock

print("Similar images by Color Moments using distance: Manhattan (cityblock)")

top_k_by_color = top_k_ranker(K, color_vector, color_tensor_dict, cityblock)

for i in range(len(top_k_by_color)):

    id = top_k_by_color[i][1]
    distance = top_k_by_color[i][0]

    print(str(i + 1) + ".", "\tImage ID: ", id, "\t\tDistance: ", distance)
    display(dataset[id][0])

```

Similar images by Color Moments using distance: Manhattan (cityblock)

1. Image ID: 2500 Distance: 0.0



2. Image ID: 2301 Distance: 10972.34



3. Image ID: 2519 Distance: 11094.308



4. Image ID: 2524 Distance: 11336.744



5. Image ID: 2308 Distance: 12073.262



6. Image ID: 2196 Distance: 12535.328



7. Image ID: 2245 Distance: 12603.198



8. Image ID: 2410 Distance: 13100.132



9. Image ID: 5007

Distance: 13128.953



10. Image ID: 2670

Distance: 13422.051



```
In [ ]: # Convert to grayscale and partition to grid.  
  
from torchvision.transforms import Grayscale  
  
gs_image = Grayscale()(image300x100)  
  
gs_grid = partition_to_grid(gs_image, num_rows=10, num_cols=10, hor_pixels=30, ver_pixels=10)
```

```
In [ ]: # Histograms of oriented gradients (HOG).  
  
# Find K most similar images by HOG.  
  
from feature_extractors.HOG import get_hog_vector  
  
hog_vector = get_hog_vector(gs_grid)  
  
# Load hog tensor database into memory.  
  
hog_tensor_dict = tensor_loader('hog_tensors.pt', os.path.abspath(""))  
  
# Correlation similarity: This measures the linear relationship between HOG feature vectors,  
# and performs well with features that have a strong linear correlation or dependency in  
# their gradient orientation distributions, taking into account directionality.  
  
from scipy.spatial.distance import correlation  
  
print("Similar images by Histogram of Oriented Gradients (HOG) using distance: Correlation")
```

```
top_k_by_hog = top_k_ranker(K, hog_vector, hog_tensor_dict, correlation)

for i in range(len(top_k_by_hog)):
    id = top_k_by_hog[i][1]
    distance = top_k_by_hog[i][0]

    print(str(i + 1) + ".", "\tImage ID: ", id, "\t\tDistance: ", distance)
    display(dataset[id][0])
```

Similar images by Histogram of Oriented Gradients (HOG) using distance: Correlation

1. Image ID: 2500 Distance: 0.0



2. Image ID: 2121 Distance: 0.0661546313127187



3. Image ID: 2308 Distance: 0.06905893472781599



4. Image ID: 2094 Distance: 0.07435084787559465



5. Image ID: 2673 Distance: 0.07991670203757606



6. Image ID: 5038 Distance: 0.08053593607760079



7. Image ID: 2167 Distance: 0.08069762336058162



8. Image ID: 2609 Distance: 0.08195859934637628



9. Image ID: 2301 Distance: 0.08204215381630864



10. Image ID: 2191 Distance: 0.08445441642668183



```
In [ ]: # Resize original image to 224x224 and convert to float tensor.
```

```
from torchvision.transforms import transforms
```

```
image224x224 = image.resize((224, 224))
```

```
img_tensor = transforms.Compose([transforms.ToTensor()])(image224x224)
```

```
In [ ]: # ResNet pre-processing.
```

```
# Find K most similar images by each of the 3 layers chosen from ResNet-50.
```

```
from feature_extractors.resnet_50 import get_resnet50_feature_vectors
```

```
avgpool_vector, layer3_vector, fc1000_vector = get_resnet50_feature_vectors(img_tensor)
```

```
# Load each tensor database into memory
```

```
avgpool_tensor_dict = tensor_loader('avgpool_tensors.pt', os.path.abspath(""))  
layer3_tensor_dict = tensor_loader('layer3_tensors.pt', os.path.abspath(""))  
fc1000_tensor_dict = tensor_loader('fc1000_tensors.pt', os.path.abspath(""))
```

```
In [ ]: # ResNet-AvgPool-1024.  
  
# Find K most similar images by ResNet-AvgPool-1024.  
  
# Cosine similarity: We use Cosine since the output of the ResNet layer is  
# normalized, reducing the importance of magnitude. It is good at discerning  
# semantic or structural similarity rather than their absolute feature values.  
# It also works well with sparse feature spaces such as those found in CNNs.  
  
from scipy.spatial.distance import cosine  
  
print("Similar images by ResNet-AvgPool-1024 using distance: Cosine")  
  
top_k_by_avgpool = top_k_ranker(K, avgpool_vector, avgpool_tensor_dict, cosine)  
  
for i in range(len(top_k_by_avgpool)):  
    id = top_k_by_avgpool[i][1]  
    distance = top_k_by_avgpool[i][0]  
  
    print(str(i + 1) + ".", "\tImage ID: ", id, "\t\tDistance: ", distance)  
    display(dataset[id][0])
```

Similar images by ResNet-AvgPool-1024 using distance: Cosine

1. Image ID: 2500 Distance: 0



2. Image ID: 2308 Distance: 0.06326800584793091



3. Image ID: 2167 Distance: 0.08668988943099976



4. Image ID: 2527 Distance: 0.09355634450912476



5. Image ID: 2254 Distance: 0.10535889863967896



6. Image ID: 2301 Distance: 0.11197084188461304



7. Image ID: 2398 Distance: 0.11582589149475098



8. Image ID: 2226 Distance: 0.11904072761535645



9. Image ID: 2412 Distance: 0.1205405592918396



10. Image ID: 2670 Distance: 0.12307322025299072



```
In [ ]: # ResNet-Layer3-1024.  
  
# Find K most similar images by ResNet-Layer3-1024.  
  
# Cosine similarity: We use Cosine since the output of the ResNet layer is  
# normalized, reducing the importance of magnitude. It is good at discerning  
# semantic or structural similarity rather than their absolute feature values.  
# It also works well with sparse feature spaces such as those found in CNNs.  
  
print("Similar images by ResNet-Layer3-1024 using distance: Cosine")  
  
top_k_by_layer3 = top_k_ranker(K, layer3_vector, layer3_tensor_dict, cosine)  
  
for i in range(len(top_k_by_layer3)):  
    id = top_k_by_layer3[i][1]  
    distance = top_k_by_layer3[i][0]  
  
    print(str(i + 1) + ".", "\tImage ID: ", id, "\t\tDistance: ", distance)  
    display(dataset[id][0])
```

```
Similar images by ResNet-Layer3-1024 using distance: Cosine  
1.      Image ID: 2500          Distance: 0
```




2. Image ID: 2519 Distance: 0.012333095073699951



3. Image ID: 2413 Distance: 0.013264119625091553



4. Image ID: 2217 Distance: 0.015680968761444092



5. Image ID: 2624 Distance: 0.016152143478393555



6. Image ID: 2569 Distance: 0.016269266605377197



7. Image ID: 2527 Distance: 0.016563475131988525



8. Image ID: 2625 Distance: 0.01666736602783203



9. Image ID: 2509 Distance: 0.01810145378112793



10. Image ID: 2411 Distance: 0.01848888397216797



```
In [ ]: # ResNet-FC-1000.  
        # Find K most similar images by ResNet-FC-1000.
```

```
# Cosine similarity: We use Cosine since the output of the ResNet layer is
# normalized, reducing the importance of magnitude. It is good at discerning
# semantic or structural similarity rather than their absolute feature values.
# It also works well with sparse feature spaces such as those found in CNNs.

print("Similar images by ResNet-FC-1000 using distance: Cosine")

top_k_by_fc1000 = top_k_ranker(K, fc1000_vector, fc1000_tensor_dict, cosine)

for i in range(len(top_k_by_fc1000)):

    id = top_k_by_fc1000[i][1]
    distance = top_k_by_fc1000[i][0]

    print(str(i + 1) + ".", "\tImage ID: ", id, "\t\tDistance: ", distance)
    display(dataset[id][0])
```

Similar images by ResNet-FC-1000 using distance: Cosine

1. Image ID: 2500 Distance: 0



2. Image ID: 2308 Distance: 0.05851387977600098



3. Image ID: 2167 Distance: 0.06570607423782349



4. Image ID: 2301 Distance: 0.07090252637863159



5. Image ID: 2203 Distance: 0.07465589046478271



6. Image ID: 2254 Distance: 0.07529681921005249



7. Image ID: 2253 Distance: 0.07565945386886597



8. Image ID: 2670 Distance: 0.07674413919448853



9. Image ID: 2527 Distance: 0.07712769508361816



10. Image ID: 2252 Distance: 0.07915341854095459

